

Text - code
Data - global
Heap - dynamic
Stack - local

Zombie - parent didn't yet call wait
Orphan - parent never called wait

ms=1000us

IPC: Pipe, FIFO, Shared Memory, Socket, Message Queue
Thread: -addressSpace +pcRegistersStack -osResources

Peterson's Solution

Thread 1	Thread 2
<pre>do { flag[0] = TRUE; turn = 1; while (flag[1] && turn==1) {}; // critical section flag[0] = FALSE; // remainder section } while (TRUE)</pre>	<pre>do { flag[1] = TRUE; turn = 0; while (flag[0] && turn==0) {}; // critical section flag[1] = FALSE; // remainder section } while (TRUE)</pre>

Disabling interrupts in a single-core CPU facilitates locks

```

boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}

```

```

int mutex;
init_lock (&mutex);

do {
    lock (&mutex);
    critical section
    unlock (&mutex);

    remainder section
} while(TRUE);

```

```

void init_lock (int *mutex)
{
    *mutex = 0;
}

void lock (int *mutex)
{
    while(TestAndSet(mutex))
        ;
}

void unlock (int *mutex)
{
    *mutex = 0;
}

```

```

int CAS(int *value, int oldval, int newval)
{
    int temp = *value;
    if (*value == oldval)
        *value = newval;
    return temp;
}

```

```

void lock (int *mutex) {
    while(CAS(mutex, 0, 1) != 0);
}

```

<pre> void mutex_init (mutex_t *lock) { lock->value = 0; list_init(&lock->wait_list); spin_lock_init(&lock->wait_lock); } void mutex_lock (mutex_t *lock) { spin_lock(&lock->wait_lock); while(TestAndSet(&lock->value)) { current->state = WAITING; list_add(&lock->wait_list, current); spin_unlock(&lock->wait_lock); schedule(); spin_lock(&lock->wait_lock); } spin_unlock(&lock->wait_lock); } void mutex_unlock (mutex_t *lock) { spin_lock(&lock->wait_lock); lock->value = 0; if (!list_empty(&lock->wait_list)) wake_up_process(&lock->wait_list); spin_unlock(&lock->wait_lock); } </pre>	<p>More reading: mutex.c in</p> <p>Thread waiting list</p> <p>To protect waiting list</p> <p>Thread state change</p> <p>Add the current thread to the waiting list</p> <p>Sleep or schedule another thread</p> <p>Someone is waiting for the lock</p> <p>Wake-up a waiting thread</p>
---	---

3

P - lock. V - signal unlock

```

P(semaphore *S) {
    S->lock->Acquire();
    S->value--;
    if(S->value < 0) {
        addQ(&S->list, P);
        S->lock->Release();
        schedule( );
        S->lock->Acquire();
    }
    S->lock->Release();
}

```

```

V(semaphore *S) {
    S->lock->Acquire();
    S->value++;
    if(S->value <= 0) {
        P = delQ(&S->list);
        wakeup(P);
    }
    S->lock->Release();
}

```

Monitor version

```
Mutex lock;
Condition full, empty;

produce (item)
{
    lock.acquire();
    while (queue.isFull())
        empty.wait(&lock);
    queue.enqueue(item);
    full.signal();
    lock.release();
}

consume()
{
    lock.acquire();
    while (queue.isEmpty())
        full.wait(&lock);
    item = queue.dequeue(item);
    empty.signal();
    lock.release();
    return item;
}
```

Semaphore version

```
Semaphore mutex = 1, full = 0,
empty = N;

produce (item)
{
    P(&empty);
    P(&mutex);
    queue.enqueue(item);
    V(&mutex);
    V(&full);
}

consume()
{
    P(&full);
    P(&mutex);
    item = queue.dequeue();
    V(&mutex);
    V(&empty);
    return item;
}
```

Reader-Writer Problem

```
semaphore mutex = 1, wrt = 1;
int readcount = 0;
```

Writer

```
do {
    P(wrt);
    write object resource
    V(wrt);
} while (TRUE);
```

Reader

```
do {
    P(mutex);
    readcount++;
    if (readcount == 1)
        P(wrt);
    V(mutex)
    read from object resource
    P(mutex);
    readcount--;
    if (readcount == 0)
        V(wrt);
    V(mutex)
} while (TRUE);
```

Deadlock conditions:

Mutual exclusion - one process per resource

No preemption - resource release must be voluntary
Hold and wait - holding some resources and waiting for more
Circular wait - circular dependency

Starvation can end, but deadlock can't

Banker's Algorithm

- General idea
 - Assume that each process's maximum resource demand is known in advance
 - $\text{Max}[i]$: process i 's maximum resource demand vector
 - **Pretend** each request is granted, then run the deadlock detection algorithm
 - If a deadlock is detected, the do not grant the request to keep the system in a **safe** state

Banker's Algorithm

1. Initialize **Avail** and **Finish** vectors
 $\text{Avail} = \text{FreeResources};$
For $i = 1, 2, \dots, n$, $\text{Finish}[i] = \text{false}$
 - **FreeResources**: resource vector
 $[\text{R1}, \text{R2}] = [0, 0]$
 - **Alloc[i]**: process i 's allocated resource vector:
 $\text{Alloc}[1] = [0, 1], \text{Alloc}[2] = [1, 0]$
2. Find an index i such that
 $\text{Finish}[i] == \text{false}$ AND
 $\text{Max}[i] - \text{Alloc}[i] \leq \text{Avail}$
If no such i exists, go to step 4
3. $\text{Avail} = \text{Avail} + \text{Alloc}[i]$, $\text{Finish}[i] = \text{true}$
Go to step 2
4. If $\text{Finish}[i] == \text{false}$, for some i , $1 \leq i \leq n$,
 - (a) then the system is in deadlock state
 - (b) if $\text{Finish}[i] == \text{false}$, then P_i is deadlocked
 - **Request[i]**: process i 's requesting vector:
 $\text{Request}[1] = [1, 0]$
 $\text{Request}[2] = [0, 0]$
 - **Max[i]**: process i 's maximum resource demand vector

Response time - time to complete a task

Wait time - total waiting

Scheduling latency - time till first run

FCFS - FIFO (min overhead)

SJF - min burst len first

SRTF - remaining time - new shorter interrupts longer (smallest waiting and interactive)

RR - round robin - each job executed for a quantum time, then rotate (interactive)

CFS - completely fair scheduler - pick min $\text{weightCpu\%} * \text{executedTime}$